



一、Plot at CartPole-v0 and Pendulum



二、Implement/adjustment of the network structure & each loss function

- CartPole-v0: Deep Q-learning with experience replay

```
12 class Net(nn.Module):
13     def __init__(self, input_size, hidden_size, output_size):
14         super().__init__()
15         self.linear1 = nn.Linear(input_size, hidden_size)
16         self.linear2 = nn.Linear(hidden_size, output_size)
17     def forward(self, x):
18         x = F.relu(self.linear1(x))
19         x = self.linear2(x)
20         return x
21
22 class Agent(object):
23     def __init__(self, **kwargs):
24         for key, value in kwargs.items():
25             setattr(self, key, value)
26         self.eval_net = Net(self.state_space_dim, 32, self.action_space_dim).to(device)
27         self.target_net = Net(self.state_space_dim, 32, self.action_space_dim).to(device)
28         self.target_net.load_state_dict(self.eval_net.state_dict())
```

在邊很單純的就是兩層 fully connect，中間夾了一個 relu activation 而已，latent vector 按照 spec 設成 32

比較特別的點是分成兩個 network，一個做主要的決定 action，一個來當作要 regression 的目標，而每過 K 個 iteration（這邊我設 50），就更新 regression 的目標成第一個 net 的 weight

如此就不會讓 regression 目標漂浮不定，比較好 train

而這邊 Loss 則是用 Mean Square Error Loss

- Pendulum-v0: DDPG

這邊則是用了叫做 Actor Critic 的技巧，Actor 負責決定動作，Critic 負責產生 Q 值

如此可以解決當 Action 是一個連續的值時，要找 maximum Q 的情況

而兩者的 model 也是很簡單，就是單純的 fully connect

Actor 的 fully connected 中間利用 relu 來連結，最後用 tanh 來輸出
Critic 中間用 relu，最後就不經過 activation function 了，直接輸出
而 Loss 一樣是用 Mean Square Error Loss

```
37 class ActorNet(nn.Module):
38     def __init__(self):
39         super(ActorNet, self).__init__()
40         self.fc = nn.Linear(3, 400)
41         self.fc2 = nn.Linear(400, 300)
42         self.mu_head = nn.Linear(300, 1)
43
44     def forward(self, s):
45         x = torch.relu(self.fc(s))
46         x = torch.relu(self.fc2(x))
47         u = torch.tanh(self.mu_head(x))
48         return u
49
50
51 class CriticNet(nn.Module):
52     def __init__(self):
53         super(CriticNet, self).__init__()
54         self.fc = nn.Linear(3, 400)
55         self.fc2 = nn.Linear(401, 300)
56         self.v_head = nn.Linear(300, 1)
57
58     def forward(self, s, a):
59         x = F.relu(self.fc(s))
60         x = F.relu(self.fc2(torch.cat([x, a], dim=1)))
61         state_value = self.v_head(x)
62         return state_value
```

三、Implement the training process of deep Q-learning

```
def learn(self):
    if (len(self.buffer)) < self.batch_size:
        return

    samples = random.sample(self.buffer, self.batch_size)
    s0, a0, r1, s1 = zip(*samples)
    s0 = torch.tensor(s0, dtype=torch.float).to(device)
    a0 = torch.tensor(a0, dtype=torch.long).view(self.batch_size, -1).to(device)
    r1 = torch.tensor(r1, dtype=torch.float).view(self.batch_size, -1).to(device)
    s1 = torch.tensor(s1, dtype=torch.float).to(device)

    y_true = r1 + self.gamma * torch.max(self.target_net(s1).detach(), dim=1)[0].view(self.batch_size, -1)
    y_pred = self.eval_net(s0).gather(1, a0)

    loss_fn = nn.MSELoss()
    loss = loss_fn(y_pred, y_true)

    self.optimizer.zero_grad()
    loss.backward()
    self.optimizer.step()
```

這邊一樣是用 TD learning 的技巧，不同的是 Q 值的計算變成了用 neural network 來取代
一開始先從歷史紀錄 (buffer) sample 出 batch_size 份量的 episodes
predict 值會是經過 self.eval_net(s0)後，取這個 action(a0)後的值
而 true value 則是利用 target_net(s1)後最大的值，加上這步的 reward
下面就把這兩個值做 Mean Square Error Loss，再 back propogate 就可以了

四、Implement of epsilon-greedy action select method

```
34     def update_iter(self):
35         self.steps += 1
36         self.epsi *= self.decay
37         if self.steps % self.copy_weight == 0:
38             self.target_net.load_state_dict(self.eval_net.state_dict())
39
40     def act(self, s0):
41         if random.random() < self.epsi:
42             a0 = random.randrange(self.action_space_dim)
43         else:
44             s0 = torch.tensor(s0, dtype=torch.float).view(1, -1).to(device)
45             a0 = torch.argmax(self.eval_net(s0)).item()
46         return a0
```

這邊有兩個 function 來完成 epsilon-greedy，一個是 update_iter，每次 iteration 完後就會 call 這個 function，來更新目前 iteration 次數，以及 epsi threshold 乘上一個 decay 值
而在實際決定動作的時候，會先 random 一個數值，如果小於 epsi threshold 的話，就會隨機挑一個 action 做，如此就可以完成 epsilon-greedy action 了

五、The mechanism of critic updating

Sample random minibatch of N transitions (s_j, a_j, r_j, s_{j+1}) from R

Set $y_i = r_i + \gamma Q'(s_{t+1}, \mu'(s_{t+1} | \theta^{\mu'})) | \theta^{Q'}$

Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i | \theta^Q))^2$

由於 critic 作用是 Q 值的產生，所以我們更新也只要考慮這部份就好了

我們會先對 buffer sample 一切 transitions 出來做更新用

第二行是在計算 target value，分別是 reward 和 target network 產生的值

第三行就是把 target value 和 evaluation network 產生的 Q 值做 Mean Square Error Loss，然後再更新就好了

要注意的是第二行在計算 target value 時，會用到 actor 的 target network 來產生 action，在實做時這部份要小心不要更新到了

六、The mechanism of actor updating

Update the actor policy using the sampled gradient:

$$\nabla_{\theta^{\mu}} \mu | s_i \approx \frac{1}{N} \sum_i \nabla_a Q(s, a | \theta^Q) |_{s=s_i, a=\mu(s_i)} \nabla_{\theta^{\mu}} \mu(s | \theta^{\mu}) | s_i$$

他的 loss 會是 sampled 的 Q 對 a 的 gradient，用 code 寫成的話會變成下面樣子

```
a_loss = -self.eval_cnet(s, self.eval_anet(s)).mean()
a_loss.backward()
```

actor 負責的是 action 的選擇，所以會希望他所選的 action 能盡量讓 critic 產生的 Q 值越大越好，而換成 loss 要越小越好的話就多加一個負號就可以了

七、How to calculate the gradients

對於 actor 的 gradient，前項 Q 對 a 的 gradient 可以看做是將 actor predict 的結果丟到 critic 後，對結果最 back propagate

後項則是 actor 產生的 action 結果做 gradient

所以整體來看可以看做是一連串下來做 back propagation 就可以了

八、How the code work

- CartPole-v0

首先先設定好 environment 以及 model 參數

下方是主程式，每次 environment 都會 reset，然後因為我是在遠端機器，所以就不處理 X server 的問題了，直接把 render 關掉

每次都會讓 agent 對這個 state 做動作，丟回 env 後就記錄下來
然後讓 agent 去 learn 他

```
107 env = gym.make('CartPole-v0')
108 params = {
109     'gamma': 0.95,
110     'epsi_high': 1,
111     'epsi_low': 0.001,
112     'decay': 0.995,
113     'lr': 0.0005,
114     'capacity': 10000,
115     'batch_size': 128,
116     'copy_weight': 50,
117     'state_space_dim': env.observation_space.shape[0],
118     'action_space_dim': env.action_space.n
119 }
120 agent = Agent(**params)
121 agent.load_weight()
122 score = []
123 mean = []
124 finish = False
```

後面就是一些更新次數或 epsilon 以及紀錄結果的 code，就不詳細解釋了

```
125 for episode in range(100):
126     s0 = env.reset()
127     total_reward = 1
128     while True:
129         # env.render()
130         a0 = agent.act(s0)
131         s1, r1, done, _ = env.step(a0)
132
133         if done:
134             r1 = -1
135
136         agent.put(s0, a0, r1, s1)
137
138         if done:
139             break
140
141         total_reward += r1
142         s0 = s1
143         agent.learn()
144     agent.update_iter()
145     # if total_reward==200 and not finish:
146     #     agent.save_weight()
147     #     print('save')
148     # if sum(score[-100:])/100 == 200:
149     #     finish = True
150     #     print('avg score reach 200')
151     score.append(total_reward)
152     print(f'Episode {episode}: score {total_reward}, avg {sum(score[-100:])/100}')
153     mean.append(sum(score[-100:])/100)
154     agent.plot(score, mean)
```


右圖是 Net 的架構以及用一個 Agent class 把 model 和訓練等 function 都包在一起

```
10 device = torch.device('cuda:2')
11 class Net(nn.Module):
12     def __init__(self, input_size, hidden_size, output_size):
13         super().__init__()
14         self.linear1 = nn.Linear(input_size, hidden_size)
15         self.linear2 = nn.Linear(hidden_size, output_size)
16     def forward(self, x):
17         x = F.relu(self.linear1(x))
18         x = self.linear2(x)
19         return x
20
21 class Agent(object):
22     def __init__(self, **kwargs):
23         for key, value in kwargs.items():
24             setattr(self, key, value)
25         self.eval_net = Net(self.state_space_dim, 32, self.action_space_dim).to(device)
26         self.target_net = Net(self.state_space_dim, 32, self.action_space_dim).to(device)
27         self.target_net.load_state_dict(self.eval_net.state_dict())
28
29         self.optimizer = optim.Adam(self.eval_net.parameters(), lr=self.lr)
30         self.buffer = []
31         self.steps = 0
32         self.epsi = self.epsi_high
33         self.training = True
```

再來是每次 iteration 結束後，會更新 step 次數以及 epsilon 的數值，另外次數如果到達一定程度的話就更新 target network 的參數

```
44 def update_iter(self):
45     self.steps += 1
46     self.epsi *= self.decay
47     if self.steps % self.copy_weight == 0:
48         self.target_net.load_state_dict(self.eval_net.state_dict())
49
50 def act(self, s0):
51     if random.random() < self.epsi:
52         a0 = random.randrange(self.action_space_dim)
53     else:
54         s0 = torch.tensor(s0, dtype=torch.float).view(1, -1).to(device)
55         a0 = torch.argmax(self.eval_net(s0)).item()
56     return a0
57
58 def put(self, *transition):
59     if len(self.buffer) == self.capacity:
60         self.buffer.pop(0)
61     self.buffer.append(transition)
```

```
62
63 def learn(self):
64     if (len(self.buffer)) < self.batch_size:
65         return
66
67     if not self.training:
68         return
69
70     samples = random.sample(self.buffer, self.batch_size)
71     s0, a0, r1, s1 = zip(*samples)
72     s0 = torch.tensor(s0, dtype=torch.float).to(device)
73     a0 = torch.tensor(a0, dtype=torch.long).view(self.batch_size, -1).to(device)
74     r1 = torch.tensor(r1, dtype=torch.float).view(self.batch_size, -1).to(device)
75     s1 = torch.tensor(s1, dtype=torch.float).to(device)
76
77     y_true = r1 + self.gamma * torch.max(self.target_net(s1).detach(), dim=1)[0].view(self.batch_size, -1)
78     y_pred = self.eval_net(s0).gather(1, a0)
79
80     loss_fn = nn.MSELoss()
81     loss = loss_fn(y_pred, y_true)
82
83     self.optimizer.zero_grad()
84     loss.backward()
85     self.optimizer.step()
```

最後是 learn，首先會先檢查 buffer 夠不夠 batch_size 的數量，再來就是 sample，然後做各種前處理，再做上面有說過的更新就好了

- Pendulum-v0

主程式如右，一樣先設定好 environment 環境，然後就開始跑 iteration
每次 iteration 都限制最多只有 200 回合，然後也是跑和存到 buffer，如果 buffer 是滿的話就更新參數

```
150 def main():
151     env = gym.make('Pendulum-v0')
152     env.seed(args.seed)
153
154     agent = Agent()
155
156     training_records = []
157     running_reward, running_q = -1000, 0
158     for i_ep in range(10000):
159         score = 0
160         state = env.reset()
161
162         for t in range(200):
163             action = agent.select_action(state)
164             state_, reward, done, _ = env.step(action)
165             score += reward
166             if args.render:
167                 env.render()
168             agent.store_transition(Transition(state, action, (reward + 8) / 8, state_))
169             state = state_
170             if agent.memory.isfull():
171                 q = agent.update()
172                 running_q = 0.99 * running_q + 0.01 * q
173
174         running_reward = running_reward * 0.9 + score * 0.1
175         training_records.append(TrainingRecord(i_ep, running_reward))
```

```
84 class Agent():
85     max_grad_norm = 0.5
86
87     def __init__(self):
88         self.training_step = 0
89         self.var = 1.
90         self.eval_cnet, self.target_cnet = CriticNet().float().to(device), CriticNet().float().to(device)
91         self.eval_anet, self.target_anet = ActorNet().float().to(device), ActorNet().float().to(device)
92         self.memory = Memory(10000)
93         self.optimizer_c = optim.Adam(self.eval_cnet.parameters(), lr=0.001)
94         self.optimizer_a = optim.Adam(self.eval_anet.parameters(), lr=0.0001)
95         self.Loss = nn.MSELoss()
96
97     def select_action(self, state):
98         state = torch.from_numpy(state).float().to(device).unsqueeze(0)
99         mu = self.eval_anet(state)
100         dist = Normal(mu, torch.tensor(self.var, dtype=torch.float).to(device))
101         action = dist.sample()
102         action.clamp(-2.0, 2.0)
103         return (action.item(),)
104
105     def save_param(self):
106         torch.save(self.eval_anet.state_dict(), 'ddpg_anet_params.pkl')
107         torch.save(self.eval_cnet.state_dict(), 'ddpg_cnet_params.pkl')
108
109     def store_transition(self, transition):
110         self.memory.update(transition)
```

上圖是 Agent，將訓練的東西都包在一起
對 Actor 和 Critic 上面的問題已經有看過了，這邊就直接省略說明
對於 Actor 和 Critic 一樣都會有 Evaluation net 和 Target net
Memory 則是用來存 episode history 的類別

selection action 裡面，為了讓 agent 能有一點變化，會是讓 actor 出來一個 mu，然後對此做一個 normal distribution 的抽樣，而其 variance 會逐步減小，最後為了不讓他衝過頭，會在對兩邊做 crop，超過或低於都直接變成最大最小值

```

112     def update(self):
113         self.training_step += 1
114
115         transitions = self.memory.sample(32)
116         s = torch.tensor([t.s for t in transitions], dtype=torch.float).to(device)
117         a = torch.tensor([t.a for t in transitions], dtype=torch.float).view(-1, 1).to(device)
118         r = torch.tensor([t.r for t in transitions], dtype=torch.float).view(-1, 1).to(device)
119         s_ = torch.tensor([t.s_ for t in transitions], dtype=torch.float).to(device)
120
121         with torch.no_grad():
122             q_target = r + args.gamma * self.target_cnet(s_, self.target_anet(s_))
123             q_eval = self.eval_cnet(s, a)
124
125         # update critic net
126         self.optimizer_c.zero_grad()
127         c_loss = self.Loss(q_eval, q_target).mean()
128         c_loss.backward()
129         nn.utils.clip_grad_norm_(self.eval_cnet.parameters(), self.max_grad_norm)
130         self.optimizer_c.step()
131
132         # update actor net
133         self.optimizer_a.zero_grad()
134         a_loss = -self.eval_cnet(s, self.eval_anet(s)).mean()
135         a_loss.backward()
136         nn.utils.clip_grad_norm_(self.eval_anet.parameters(), self.max_grad_norm)
137         self.optimizer_a.step()
138
139         if self.training_step % 300 == 0:
140             self.target_cnet.load_state_dict(self.eval_cnet.state_dict())
141         if self.training_step % 301 == 0:
142             self.target_anet.load_state_dict(self.eval_anet.state_dict())
143
144         self.var = max(self.var * 0.999, 0.01)
145
146         return q_eval.mean().item()

```

最後是整個 model 的更新，我們會從 buffer 裡 sample batch_size 大小的份數，然後做完各種前處理之後就是實做上面解釋過的，對 actor 和 critic 的更新了
而這邊為了方便，沒有讓 target network 做 soft update，而是跟 DQN 一樣到一定 episode 就覆蓋過去，結果也顯示還是不錯的

右圖則是存 history 的 class
更新方式是用一個 data_pointer 紀錄這串 list 的頭是哪裡，滿了後要更新就把他蓋過去

sample 也就直接用 numpy 的 random choice 來取出一定數量就完成了

```

65     class Memory():
66         data_pointer = 0
67         isfull = False
68
69         def __init__(self, capacity):
70             self.memory = np.empty(capacity, dtype=object)
71             self.capacity = capacity
72
73         def update(self, transition):
74             self.memory[self.data_pointer] = transition
75             self.data_pointer += 1
76             if self.data_pointer == self.capacity:
77                 self.data_pointer = 0
78                 self.isfull = True
79
80         def sample(self, batch_size):
81             return np.random.choice(self.memory, batch_size)

```


九、Other study or improvement for the project

其實老師在上這邊時非常的快速，我是自己額外看完 youtube 上李宏毅影片後才寫這份作業的，但其實自己寫的部份也沒有說非常多，在網路上有找到別人寫好的版本，參考後修改一些地方就下去跑了


其中也發現別人包成 class 的寫法和對於 sample history data 的寫法，甚至是 model 參數的設定方式都不一樣，收穫十分良多

pendulum 會要跑好幾個小時，但幸好沒太多問題，很快就 train 起來了

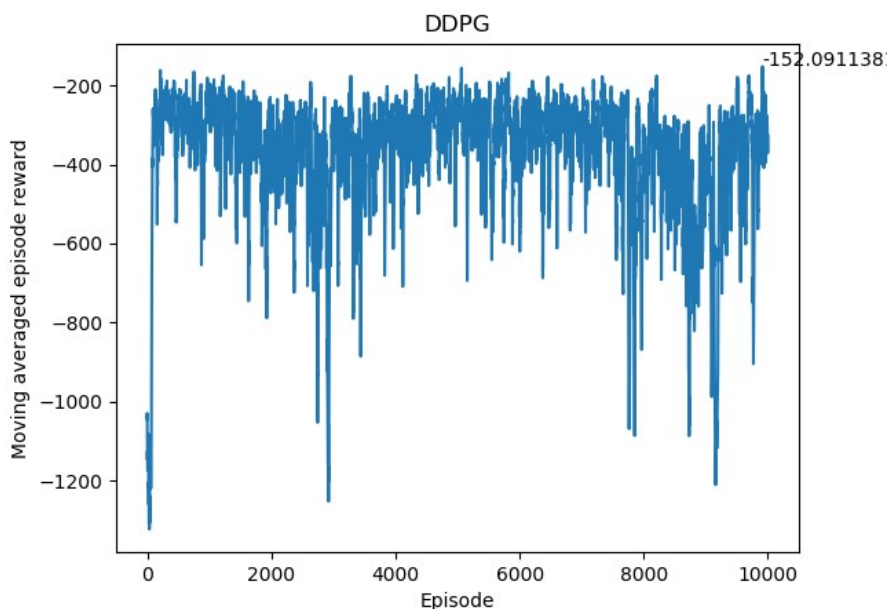
cartpole 則是 train 到很好的時候再多 train 就會不小心壞掉，應該是跟那個 epsilon 有關係，所以我在 training 階段，只要連續到 200 一定次數就不再 store weight，結果就可以像下面一樣都到 200 了

十、Performance

- CartPole: 200



- Pendulum: -152.091138



```
Episode 66: score 200.0, avg 134.0
Episode 67: score 200.0, avg 136.0
Episode 68: score 200.0, avg 138.0
Episode 69: score 200.0, avg 140.0
Episode 70: score 200.0, avg 142.0
Episode 71: score 200.0, avg 144.0
Episode 72: score 200.0, avg 146.0
Episode 73: score 200.0, avg 148.0
Episode 74: score 200.0, avg 150.0
Episode 75: score 200.0, avg 152.0
Episode 76: score 200.0, avg 154.0
Episode 77: score 200.0, avg 156.0
Episode 78: score 200.0, avg 158.0
Episode 79: score 200.0, avg 160.0
Episode 80: score 200.0, avg 162.0
Episode 81: score 200.0, avg 164.0
Episode 82: score 200.0, avg 166.0
Episode 83: score 200.0, avg 168.0
Episode 84: score 200.0, avg 170.0
Episode 85: score 200.0, avg 172.0
Episode 86: score 200.0, avg 174.0
Episode 87: score 200.0, avg 176.0
Episode 88: score 200.0, avg 178.0
Episode 89: score 200.0, avg 180.0
Episode 90: score 200.0, avg 182.0
Episode 91: score 200.0, avg 184.0
Episode 92: score 200.0, avg 186.0
Episode 93: score 200.0, avg 188.0
Episode 94: score 200.0, avg 190.0
Episode 95: score 200.0, avg 192.0
Episode 96: score 200.0, avg 194.0
Episode 97: score 200.0, avg 196.0
Episode 98: score 200.0, avg 198.0
Episode 99: score 200.0, avg 200.0
(kk) karljackab@pc3421:~/DL/lab8$
```